


AJAR — HCDS methodology pipeline

Are LLMs Just Acting Reasonable? Hidden CoT Detection in Qwen3-4B (Instruct vs Thinking)

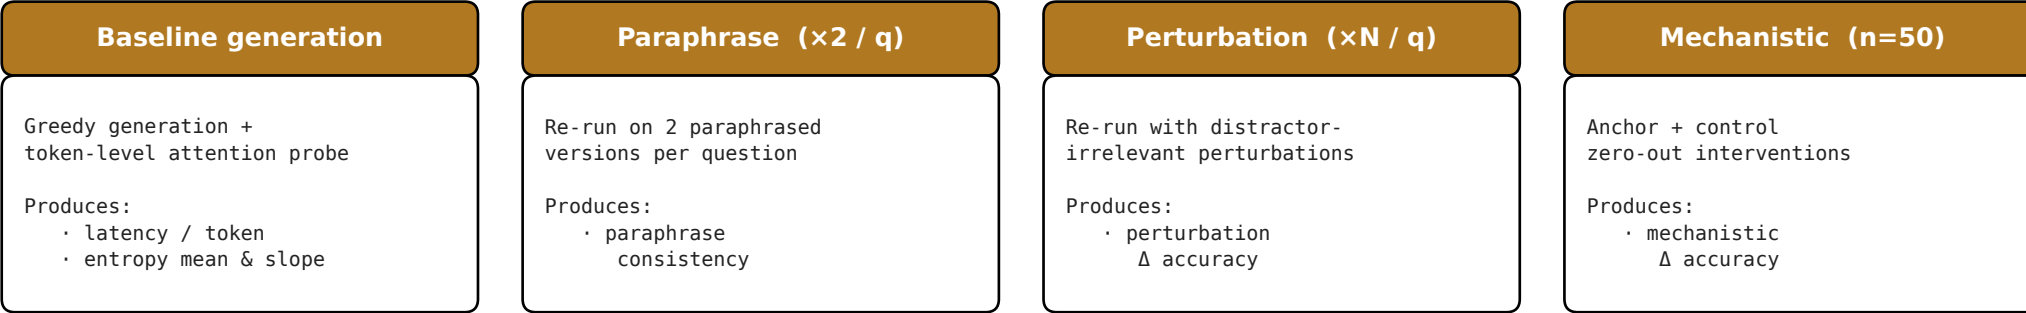
– SETUP



run $2 \times 3 \times n = 300$ cells ($n=50$) / 3000 cells ($n=500$)

– PIPELINE

For each cell, four parallel sub-runs each produce a slice of the feature vector:



z-score per model across all (question × prompt) rows → feature vector $f(q, \text{prompt}) \in \mathbb{R}^6$ for every cell

– RESULT

