



Output length distribution by prompt condition (GSM8K, n=500)

Qwen3-4B Instruct



Qwen3-4B Thinking



Instruct + explicit_no_cot $\approx$ 6 median tokens (true answer-only). Thinking + explicit_no_cot $\approx$ 491 median — model ignores the no-CoT instruction.